



BST & Tree Traversals

The Complete Deep-Dive Guide

Insertion · Deletion · Search · Traversals · Recursion Explained Step-by-Step

☐ **BST Operations**

☐ **4 Traversals**

☐ **Recursion Explained**

Table of Contents

Table of Contents	2
1.1 From Binary Tree to BST — The One Rule That Changes Everything	4
1.2 Height and Time Complexity — The Core Insight.....	4
2.1 Node Structure	6
Python Implementation	6
C / C++ Implementation	6
3.1 The Insertion Logic.....	8
3.2 Step-by-Step: Building the BST	8
3.3 Recursive Code — Explained Line by Line	9
Python	9
C++ Implementation	10
3.4 Recursion Call Stack — Full Trace for insert(20→15→30→12).....	11
4.1 How Search Works.....	13
4.2 Search Code — Every Line Explained.....	13
Python	13
C++	14
4.3 Recursion Trace for search(root=20, value=18).....	14
5.1 Predecessor and Successor — Must Know First	16
5.2 The 3 Cases of Deletion	17
Case 1: Deleting a LEAF Node (no children)	17
Case 2: Deleting a Node with ONE child.....	17
Case 3: Deleting a Node with TWO children (the tricky one!).....	17
5.3 In-Order Successor Helper Function	18
Python	18
5.4 Delete Function — Full Code with Every Line Explained	19
Python	19
C++	20
5.5 Deletion Recursion Trace — Deleting 12 (Leaf).....	21
6.1 The Traversal Problem	23
7.1 The Pattern	24
7.2 Code — Line by Line	24
Python	24
C++	25
7.3 Full Recursion Trace — Step by Step.....	25
8.1 The Pattern + Secret Superpower	27
8.2 Code — Line by Line	27

Python	27
C++	28
8.3 Full Recursion Trace	28
9.1 The Pattern	30
9.2 Code — Line by Line	30
Python	30
C++	31
9.3 Full Recursion Trace	31
10.1 Full Python Program.....	33
10.2 Full C++ Program	35
11.1 BST Operations.....	37
11.2 Traversal Complexities	37
11.3 Key Rules to Memorize	38

1. Binary Search Tree — The Big Picture

Why BST exists and what makes it special

1.1 From Binary Tree to BST — The One Rule That Changes Everything

A Binary Tree is just any tree where each node has at most 2 children. It has NO ordering rule. So if you want to find a value, you might have to check every single node — $O(n)$ in the worst case.

A Binary Search Tree adds ONE powerful rule:


□ The BST Rule:

LEFT subtree → all values are SMALLER than the parent node

RIGHT subtree → all values are LARGER than the parent node

This applies recursively to EVERY node in the tree — not just the root!

Binary Tree (no rule):



To find 9 in Binary Tree:

Check 5 → not it

Check 8 → not it


Check 2 → not it

Check 1 → not it

Check 9 → FOUND! (5 checks)

Binary Tree search: $O(n)$

Binary SEARCH Tree (BST):



To find 18 in BST:

Check 20 → $18 < 20$ → go LEFT

Check 15 → $18 > 15$ → go RIGHT

Check 18 → FOUND! ✓


Only 3 checks vs up to n checks!

BST search: $O(\log n)$ ✓

1.2 Height and Time Complexity — The Core Insight

The number of comparisons needed to find any element = the HEIGHT of the tree. And the height of a balanced BST with n nodes = $\log_2(n)$.

n=7 nodes, balanced BST:



Height = $\log_2(7) \approx 3$

Level 0: 1 node

Level 1: 2 nodes

Level 2: 4 nodes

Worst case: 3 comparisons to find ANY element in 7-node tree!
 Compare to linear search: up to 7 comparisons.
 For $n=1,000,000$ nodes: BST needs only $\log_2(1M) \approx 20$ comparisons!

$n=1$ → height=0 → 1 comparison max
 $n=3$ → height=1 → 2 comparisons max
 $n=7$ → height=2 → 3 comparisons max
 $n=15$ → height=3 → 4 comparisons max
 $n=1000000$ → height=20 → 20 comparisons max ← incredibly efficient!

Metric	Binary Tree	Binary Search Tree (balanced)	Why it matters
Search	$O(n)$	$O(\log n)$	BST is exponentially faster for large datasets
Insert	$O(1)$ at known pos	$O(\log n)$	BST must find the right spot first
Delete	$O(n)$	$O(\log n)$	BST can navigate directly to node
Sorted output	$O(n \log n)$ sort needed	$O(n)$ inorder	BST gives sorted data FREE

□ Interview Trap — 3 Questions That Are The Same!

Q1: Print all elements of a BST in increasing order

Q2: How to check if a tree is a valid BST

Q3: Perform inorder traversal of a BST

Answer: ALL THREE reduce to INORDER TRAVERSAL (Left → Root → Right).

Interviewers phrase it differently to confuse you. Now you know!

2. The Node — Building Block of BST

What each node stores and why

2.1 Node Structure

Every BST is made of nodes. Each node stores three things: the data value, a pointer to the left child, and a pointer to the right child.

```
+-----+-----+-----+
|  LEFT  |  DATA  |  RIGHT  |
| pointer | value   | pointer |
+-----+-----+-----+
      |               |
      v               v
  (smaller          (larger
   values)          values)
```

Python Implementation

```
class Node:
    def __init__(self, value):
        # The actual data this node holds
        self.data = value

        # Pointer to LEFT child (smaller values live here)
        # None means 'no left child' (this node is a leaf on left side)
        self.left = None

        # Pointer to RIGHT child (larger values live here)
        # None means 'no right child'
        self.right = None

# Time: O(1) — just initializing 3 attributes
# Space: O(1) — only 3 fields stored per node
```

C / C++ Implementation

```
// C struct approach
struct Node {
    int data;           // value stored
    Node* left;         // pointer to left child
    Node* right;        // pointer to right child
};

// C++ class approach (constructor)
class Node {
```

```
public:
    int data;
    Node* left;
    Node* right;

    Node(int value) {
        data = value;
        left = nullptr;    // nullptr = no child
        right = nullptr;
    }
};
// Time: O(1), Space: O(1)
```

+ 3. Insertion in BST

How values are placed — the recursive journey

3.1 The Insertion Logic

To insert a value into a BST, we follow the BST rule at every node:

- If the tree is empty → the new node becomes the root
- If value < current node → go LEFT
- If value > current node → go RIGHT
- If value == current node → duplicate, ignore it (BSTs usually don't store duplicates)
- Keep going until we find an empty (None/null) spot → insert there

3.2 Step-by-Step: Building the BST

Let's insert: 20, 15, 30, 12, 18, 40 — one by one.

```

Step 1: Insert 20
Tree is empty → 20 becomes root
[20]

Step 2: Insert 15
15 < 20 → go LEFT of 20
20's left is None → place 15 there
  [20]
  /
 [15]

Step 3: Insert 30
30 > 20 → go RIGHT of 20
20's right is None → place 30 there
  [20]
  /  \
 [15] [30]

Step 4: Insert 12
12 < 20 → go LEFT (reach 15)
12 < 15 → go LEFT (reach None)
Place 12 at 15's left
  [20]
  /  \
 [15] [30]
 /
[12]

Step 5: Insert 18

```



```

18 < 20 → go LEFT (reach 15)
18 > 15 → go RIGHT (reach None)
Place 18 at 15's right
    [20]
    /  \
   [15] [30]
  /  \
 [12] [18]

Step 6: Insert 40
40 > 20 → go RIGHT (reach 30)
40 > 30 → go RIGHT (reach None)
Place 40 at 30's right
    [20]
    /  \
   [15] [30]
  /  \   \
 [12] [18] [40] ← FINAL TREE

```

3.3 Recursive Code — Explained Line by Line

Python

```

def insertion_in_bst(root, value):
    # ——— BASE CASE ———
    # We've reached an empty spot in the tree (None node).
    # This is exactly where the new value belongs!
    # Create a new Node here and return it so the parent can
    # connect to it via root.left = ... or root.right = ...
    #
    # Time: O(1) for this operation
    # Space: O(1) for the new node created
    if root is None:
        return Node(value)          # new node born here

    # ——— DUPLICATE CHECK ———
    # BSTs typically don't allow duplicates.
    # If value already exists, just return root unchanged.
    elif value == root.data:
        return root                 # ignore duplicate

    # ——— GO LEFT ———
    # value is smaller → it belongs in the LEFT subtree.
    # We recursively insert into the left subtree and
    # REASSIGN root.left to the result.
    # Why reassign? Because the recursive call might return a
    # brand new Node (when root.left was None) — we must link it!
    elif value < root.data:
        root.left = insertion_in_bst(root.left, value)

```

```
# — GO RIGHT —————
# value is larger → it belongs in the RIGHT subtree.
# Same logic as left, but for right side.
else:
    root.right = insertion_in_bst(root.right, value)

# — RETURN ROOT —————
# CRITICAL: always return root so the parent call can
# reconnect the subtree correctly.
# Without this, the tree gets disconnected!
return root

# Time Complexity:
# Best:    O(1)      - tree is empty, just create root
# Average: O(log n)  - balanced tree, traverse half height
# Worst:   O(n)      - skewed tree (sorted input), traverse all nodes
# Space Complexity:
# O(h) - h = height of tree (recursion call stack frames)
# Balanced: O(log n)   Skewed: O(n)
```

C++ Implementation

```
Node* insertBST(Node* root, int value) {
    // Base case: empty spot found → create new node here
    // nullptr in C++ = None in Python
    if (root == nullptr)
        return new Node(value);

    // Duplicate: ignore
    if (value == root->data)
        return root;

    // Go LEFT: value smaller than current node
    // root->left uses -> because root is a POINTER
    if (value < root->data)
        root->left = insertBST(root->left, value);

    // Go RIGHT: value larger than current node
    else
        root->right = insertBST(root->right, value);

    // Return root to maintain parent linkage
    return root;
}

// Usage:
// Node* root = nullptr;
// root = insertBST(root, 20);
// root = insertBST(root, 15);
// root = insertBST(root, 30);
```

3.4 Recursion Call Stack — Full Trace for insert(20→15→30→12)

Let's trace EXACTLY what happens in memory when we insert 12 into the tree that already has 20, 15, 30:

Call: insertion_in_bst(root=Node(20), value=12)

```
Frame 1: root=Node(20), value=12
12 < 20 → go LEFT
root.left = insertion_in_bst(Node(15), 12) ← RECURSE
(waiting for recursive call to return...)
```

↓ recursive call

```
Frame 2: root=Node(15), value=12
12 < 15 → go LEFT
root.left = insertion_in_bst(None, 12) ← RECURSE
(waiting for recursive call to return...)
```

↓ recursive call

```
Frame 3: root=None, value=12
BASE CASE HIT! root is None
return Node(12) ← NEW NODE CREATED
```

↑ returns Node(12)

```
Frame 2 resumes: root.left = Node(12)
Node(15).left is now Node(12) ✓
return Node(15) ← return current root
```

↑ returns Node(15)

```
Frame 1 resumes: root.left = Node(15)
Node(20).left is still Node(15) ✓ (unchanged, relinked)
return Node(20) ← final return
```

Final tree:

```

      [20]
     /  \
    [15] [30]
   /
  [12] ← successfully inserted!
```

Case	Scenario	Time	Space	Example
Best	Tree is empty — new node becomes root	$O(1)$	$O(1)$	Insert into empty tree
Average	Balanced BST — traverse $\sim \log(n)$ nodes	$O(\log n)$	$O(\log n)$	Insert 25 into balanced 7-node tree
Worst	Skewed tree — traverse all n nodes	$O(n)$	$O(n)$	Insert 1,2,3,4,5 in order → right skewed


□ 4. Search in BST

The power of $O(\log n)$ — finding elements fast

4.1 How Search Works

Searching in a BST is like playing a 'higher or lower' guessing game. At each node, you know exactly which direction to go — left (if target is smaller) or right (if target is larger). This eliminates half the remaining tree at every step.

BST for search examples:



Search for 18:

```

Step 1: Compare 18 with root 20 → 18 < 20 → go LEFT
Step 2: Compare 18 with node 15 → 18 > 15 → go RIGHT
Step 3: Compare 18 with node 18 → 18 == 18 → FOUND! ✓
Only 3 comparisons for 6-node tree!
  
```

Search for 25 (not in tree):

```

Step 1: Compare 25 with root 20 → 25 > 20 → go RIGHT
Step 2: Compare 25 with node 30 → 25 < 30 → go LEFT
Step 3: 30's left is None → NOT FOUND
3 comparisons to confirm 25 is absent!
  
```

4.2 Search Code — Every Line Explained

Python

```

def search_in_bst(root, value):
    # — BASE CASE 1: Not found —————
    # We've gone past a leaf node (reached None/null).
    # The value doesn't exist anywhere in this subtree.
    # This happens when we kept going left or right and fell off the tree.
    if root is None:
        print(f'Element {value} not found')
        return False

    # — BASE CASE 2: Found! —————
    # Current node's data matches what we're looking for.
    # No need to go further — report success.
    elif value == root.data:
        print(f'Element {value} found')
  
```

```

        return True

# ——— RECURSIVE CASE: Go LEFT ———
# Target value is SMALLER than current node.
# By BST property, it can ONLY exist in the left subtree.
# Right subtree guaranteed to not have it → skip entirely!
elif value < root.data:
    return search_in_bst(root.left, value)

# ——— RECURSIVE CASE: Go RIGHT ———
# Target value is LARGER than current node.
# By BST property, it can ONLY exist in the right subtree.
else:
    return search_in_bst(root.right, value)

# Time Complexity:
#   Best:    O(1)      - value is at root
#   Average: O(log n)  - balanced BST, cut tree in half each step
#   Worst:   O(n)      - skewed tree OR value at deepest leaf
# Space Complexity:
#   O(h) due to recursion stack (h = height of tree)
#   Balanced: O(log n)   Skewed: O(n)

```

C++

```

bool searchBST(Node* root, int value) {
    // Base case 1: fell off the tree → not found
    if (root == nullptr)
        return false;

    // Base case 2: exact match → found!
    if (value == root->data)
        return true;

    // Recursive: search left or right based on comparison
    if (value < root->data)
        return searchBST(root->left, value);
    else
        return searchBST(root->right, value);
}

```

4.3 Recursion Trace for search(root=20, value=18)

Call Stack builds up (top = most recent call):

```

| search(root=Node(18), value=18) |
| 18 == 18 → BASE CASE 2: return True ← unwinds here |
|_____|

```

```
| search(root=Node(15), value=18) |  
| 18 > 15 → go RIGHT → call search(Node(18), 18) |  
| returns True ← gets result, passes it up |  
|-----|  
| search(root=Node(20), value=18) ← FIRST call |  
| 18 < 20 → go LEFT → call search(Node(15), 18) |  
| returns True ← final answer bubbles up to caller |
```

Each frame is pushed when we recurse, popped when we return.
Memory freed for each frame as soon as it returns.

□ 5. Deletion in BST

The hardest operation — 3 cases, explained clearly

5.1 Predecessor and Successor — Must Know First

Before deletion, understand two key concepts:

□ In-Order Predecessor and Successor:

Given a node N in a BST after inorder traversal [6,8,9,10,20,25,30,40,50]:

In-Order Predecessor of N = the element just BEFORE N in sorted order

In-Order Successor of N = the element just AFTER N in sorted order

For N=25: Predecessor=20 (one step LEFT then rightmost),


Successor=30 (one step RIGHT then leftmost)

For N=10: Predecessor=9, Successor=20

How to find In-Order SUCCESSOR of node N:

1. Go ONE step to the RIGHT child of N
2. Then go as FAR LEFT as possible
3. The leftmost node = smallest value larger than N = Successor

Example: Successor of 25 in tree below:



Step 1: from 25, go RIGHT → reach 28... wait, 25's right = 28

Actually 25 has no right child shown, go to parent's right subtree...

Cleaner example — Successor of 10:

Step 1: Go right from 10 → reach 15

Step 2: Go as far left as possible from 15 → 15 has no left child

Result: Successor of 10 = 15 ✓

How to find In-Order PREDECESSOR of node N:

1. Go ONE step to the LEFT child of N

2. Then go as FAR RIGHT as possible
3. The rightmost node = largest value smaller than N = Predecessor

5.2 The 3 Cases of Deletion

When deleting a node from a BST, there are exactly 3 situations. Each is handled differently.

Case 1: Deleting a LEAF Node (no children)

```

Delete node 12 (a leaf):          After deletion:
    [20]                          [20]
   /  \                          /  \
  [15] [30]                      [15] [30]
 /  \  \                          \  \
[12] [18] [40]                  [18] [40]
  
```

Action: Simply return None to the parent.

Parent (15) had root.left = Node(12).

After deletion, root.left = None.

Memory for Node(12) is freed.

Code handles it automatically: if root.left is None → return root.right

When BOTH are None (leaf), returns None → parent disconnects it.

Case 2: Deleting a Node with ONE child

```

Delete node 30 (has only right child 40):
    [20]                          [20]
   /  \                          /  \
  [15] [30]                      [15] [40]
 /  \  \                          /  \
[12] [18] [40]                  [12] [18]
  
```

Action: Connect 30's parent directly to 30's child.

Parent (20) had root.right = Node(30).

After deletion, root.right = Node(40). (bypasses 30)

Node(30) is removed. Node(40) takes its place.

Code: if root.left is None → return root.right (right child takes over)

if root.right is None → return root.left (left child takes over)

Case 3: Deleting a Node with TWO children (the tricky one!)

Delete node 20 (has both left=15 and right=30):

PROBLEM: Can't just remove it — two subtrees need a new root!

SOLUTION: Replace with IN-ORDER SUCCESSOR (or predecessor)

Why successor? Because it is:

- The smallest value LARGER than 20
- Guaranteed to be > everything in left subtree (maintains BST!)
- Guaranteed to be < everything in right subtree (maintains BST!)

Full BST for this example:



Steps:

1. Find inorder successor of 20:
Go RIGHT to 30, then go as FAR LEFT as possible → reach 25
Successor = 25
2. COPY successor's value (25) into the node being deleted (20)
Tree now has two 25s: one at old root, one at original position
[6,8,9,10,15,18,25,25,30,40,50] – still valid BST!
3. DELETE the original successor node (25) from right subtree
25 has no left child → Case 2 applies (easy delete)

Result: Node that had 20 now has 25, and old 25 is gone. ✓

5.3 In-Order Successor Helper Function

Python

```

def get_in_order_successor(root):
    """
    Find the in-order successor of a node.
    Strategy: go one step RIGHT, then as far LEFT as possible.
    Result: smallest value in the right subtree = in-order successor.
    """
    # Step 1: Move to the right child of the node to be deleted
    # This right child is larger than the deleted node
    root = root.right

    # Step 2: Keep going left until we can't anymore
    # The leftmost node of the right subtree = smallest value
    # that is still larger than the deleted node
    while root is not None and root.left is not None:
        root = root.left

    # Return the successor node (not just its value, the whole node)
    return root
  
```

```
# Example trace for finding successor of 20 in tree above:
# root.right = Node(30)    ← step right
# Node(30).left = Node(25) ← step left
# Node(25).left = None     ← stop, leftmost reached
# Return Node(25) → 25 is the successor of 20 ✓

# Time: O(h) — traverse from right child to leftmost leaf
# Space: O(1) — just moves a pointer, no recursion
```

5.4 Delete Function — Full Code with Every Line Explained

Python

```
def delete_from_bst(root, value):
    # — BASE CASE: Value not in tree —————
    # Reached None → value was never found → nothing to delete
    if root is None:
        print(f'Element {value} not found')
        return None

    # — NAVIGATE LEFT —————
    # Value is smaller → the node to delete is in the left subtree
    # Recurse left and REASSIGN to update the link correctly
    elif value < root.data:
        root.left = delete_from_bst(root.left, value)

    # — NAVIGATE RIGHT —————
    # Value is larger → the node to delete is in the right subtree
    elif value > root.data:
        root.right = delete_from_bst(root.right, value)

    # — FOUND THE NODE TO DELETE —————
    # root.data == value → this is our target!
    else:
        # Case 1 & 2: Only right child (or no children at all)
        # Return root.right → parent will link to it directly
        # If right is also None (leaf), returns None → parent disconnects
        if root.left is None:
            return root.right

        # Case 2 (mirror): Only left child
        # Return root.left → parent links to left child
        elif root.right is None:
            return root.left

        # Case 3: Two children — the complex case
        # Cannot remove directly: need to maintain BST structure
        # Strategy: find inorder successor, copy its value here,
```

```

        # then delete the original successor from right subtree
        succ = get_in_order_successor(root)
        # Copy successor's value into this node
        root.data = succ.data
        # Delete successor from right subtree
        # Successor has at most ONE child (no left child by definition!)
        # So this deletion will fall into Case 1 or Case 2
        root.right = delete_from_bst(root.right, succ.data)

# ——— RETURN ROOT ———
# Always return root so parent links stay intact
return root

# Time Complexity:
#   Best:    O(1)      — root is the node to delete, it's a leaf
#   Average: O(log n)  — balanced BST
#   Worst:   O(n)      — skewed tree
# Space Complexity: O(h) — recursion call stack

```

C++

```

Node* getSuccessor(Node* root) {
    root = root->right;           // one step right
    while (root && root->left)     // go as far left as possible
        root = root->left;
    return root;
}

Node* deleteBST(Node* root, int value) {
    if (root == nullptr) return nullptr; // not found

    if (value < root->data)
        root->left = deleteBST(root->left, value);
    else if (value > root->data)
        root->right = deleteBST(root->right, value);
    else {
        // Case 1 & 2: 0 or 1 child
        if (root->left == nullptr) {
            Node* temp = root->right;
            delete root;           // free memory in C++!
            return temp;
        }
        if (root->right == nullptr) {
            Node* temp = root->left;
            delete root;
            return temp;
        }
        // Case 3: two children
        Node* succ = getSuccessor(root);
        root->data = succ->data;
    }
}

```

```

    root->right = deleteBST(root->right, succ->data);
}
return root;
}
// Note: In C++, always 'delete' the node to free heap memory!
// Python handles this automatically (garbage collection).

```

5.5 Deletion Recursion Trace — Deleting 12 (Leaf)

Tree before: [20]->[15]->[12] (left), [18] (right), [30]->[40]
 Call: delete_from_bst(root=Node(20), value=12)

```

Frame 1: root=Node(20), value=12
12 < 20 → go LEFT
root.left = delete_from_bst(Node(15), 12) ← RECURSE
(waiting...)

```

↓

```

Frame 2: root=Node(15), value=12
12 < 15 → go LEFT
root.left = delete_from_bst(Node(12), 12) ← RECURSE
(waiting...)

```

↓

```

Frame 3: root=Node(12), value=12
12 == 12 → FOUND! Target node!
root.left is None AND root.right is None → LEAF
if root.left is None → return root.right = None

```

↑ returns None

```

Frame 2 resumes: root.left = None
Node(15).left is now None (12 is gone) ✓
return Node(15)

```

↑ returns Node(15)

```

Frame 1 resumes: root.left = Node(15) (unchanged, relinked)
return Node(20)

```

Tree after: [20]->[15]->[18] (right only), [30]->[40]
 Node 12 is gone. Tree is still a valid BST. ✓

Case	Scenario	Action	Time	Space
Case 1	Leaf node (0 children)	Return None to parent. Parent disconnects it.	$O(\log n)$ avg	$O(h)$
Case 2	1 child	Return the child to parent. Child takes deleted node's place.	$O(\log n)$ avg	$O(h)$
Case 3	2 children	Find successor, copy value, delete successor (Case 1 or 2).	$O(\log n)$ avg	$O(h)$

□ 6. Tree Traversals — Why We Need Them

A tree isn't linear — so how do we visit every node?

6.1 The Traversal Problem

In an array or linked list, visiting every element is trivial — just go from left to right. But a tree branches! From node 1, you can go to node 2 OR node 3. Which do you visit first? This is the traversal question.

```
Simple tree:      [1]
                  /  \
                 [2]  [3]
```

All possible orderings of {1, 2, 3}:

```
1 2 3 | 1 3 2 | 2 1 3 | 2 3 1 | 3 1 2 | 3 2 1
(that is 3! = 6 possible orders)
```

We focus on ROOT position, always going LEFT before RIGHT:

Category 1 — Root at START (Preorder):

```
Root → Left → Right → 1 2 3
```

Category 2 — Root in MIDDLE (Inorder):

```
Left → Root → Right → 2 1 3
```

Category 3 — Root at END (Postorder):

```
Left → Right → Root → 2 3 1
```

3 standard traversals selected (one from each category).

(We skip Root→Right→Left etc. by convention: always left before right)

□ Why These 3 Specifically?

By convention: we ALWAYS visit the LEFT child before the RIGHT child.

This halves the choices: 6 orderings → 3 standard traversals.

Each has a specific use case — not arbitrary choices!

1 7. Pre-Order Traversal

Root → Left → Right

7.1 The Pattern

Visit the ROOT first, then recursively traverse the LEFT subtree, then the RIGHT subtree. You see the parent BEFORE its children.

```

Tree:                Preorder Result:
    [1]
   /  \
  [3]  [5]
 /  \  \
[2] [4] [8]

Visit root → print 1
Go left   → print 3
3's left  → print 2
3's right → print 4
Go right  → print 5
5's right → print 8

Preorder: 1 → 3 → 2 → 4 → 5 → 8

```

7.2 Code — Line by Line

Python

```

def pre_order_traversal(root):
    # — BASE CASE —————
    # Empty node → nothing to visit → return empty list
    # This triggers when we try to go left/right from a leaf
    #
    # Time: O(1) for this check
    if root is None:
        return []

    # — VISIT ROOT FIRST (Pre = before children) —————
    # Put current node's data in result list BEFORE recursing
    # This is what makes it 'pre'-order
    #
    # Then recurse left: get all data from left subtree as a list
    # Then recurse right: get all data from right subtree as a list
    # Concatenate: [root] + [left subtree] + [right subtree]
    return (
        [root.data]                + # root first
        pre_order_traversal(root.left) + # then left subtree
        pre_order_traversal(root.right)  # then right subtree
    )

# Time Complexity:

```



```
# Best/Average/Worst: O(n) — must visit every single node
# Space Complexity:
# O(h) for call stack + O(n) for the result list
# Balanced: O(log n) stack depth   Skewed: O(n) stack depth
```

C++

```
void preOrder(Node* root, vector<int>& result) {
    if (root == nullptr) return; // base case

    result.push_back(root->data); // visit root FIRST
    preOrder(root->left, result); // then left subtree
    preOrder(root->right, result); // then right subtree
}

// Alternative: return vector (cleaner but slightly less efficient)
vector<int> preOrder(Node* root) {
    if (!root) return {};
    vector<int> res = {root->data};
    auto left = preOrder(root->left);
    auto right = preOrder(root->right);
    res.insert(res.end(), left.begin(), left.end());
    res.insert(res.end(), right.begin(), right.end());
    return res;
}
```

7.3 Full Recursion Trace — Step by Step

Tree: Node(1), left=Node(3)[left=Node(2), right=Node(4)], right=Node(5)[right=Node(8)]

```
pre_order_traversal(Node(1)) ← initial call
├─ root=Node(1), not None
├─ [1] + pre_order(Node(3)) + pre_order(Node(5))
├─
├─ pre_order(Node(3)) :
│   └─ [3] + pre_order(Node(2)) + pre_order(Node(4))
│   └─
│   └─ pre_order(Node(2)) :
│       └─ [2] + pre_order(None) + pre_order(None)
│       └─ pre_order(None) → [] ← base case
│       └─ pre_order(None) → [] ← base case
│       └─ returns [2]
│   └─ pre_order(Node(4)) :
│       └─ [4] + pre_order(None) + pre_order(None)
│       └─ returns [4]
└─ returns [3] + [2] + [4] = [3, 2, 4]
```

```

└─ pre_order(Node(5)) :
    └─ [5] + pre_order(None) + pre_order(Node(8))
        └─ pre_order(Node(8)) :
            └─ returns [8]
        └─ returns [5] + [] + [8] = [5, 8]
└─ Final: [1] + [3,2,4] + [5,8] = [1, 3, 2, 4, 5, 8] ✓

```

□ Preorder Use Cases:

Copying/cloning a tree: insert nodes in preorder into new tree = exact copy

Serializing to file: store preorder sequence → rebuild identical tree later

Prefix expression evaluation: math like + * 3 4 2 uses preorder notation

Directory listing: print folder name before its contents (ls -R style)

2 8. In-Order Traversal

Left → Root → Right = Sorted Output!

8.1 The Pattern + Secret Superpower

Visit the LEFT subtree first, then the ROOT, then the RIGHT subtree. For a BST, this ALWAYS produces output in sorted ascending order. This is the most important traversal for BSTs.

```
BST:           Inorder Result:
      [1]
     /  \
    [3]  [5]
   /  \  \
  [2] [4] [8]
      Go far left  → visit 2
      Back to 3    → visit 3
      3's right    → visit 4
      Back to root → visit 1
      Go right     → visit 5
      5's right    → visit 8
```

Inorder: 2 → 3 → 4 → 1 → 5 → 8

For a BST with values [12,15,18,20,30,40]:

Inorder always gives: 12 → 15 → 18 → 20 → 30 → 40 (sorted!) ✓

8.2 Code — Line by Line

Python

```
def in_order_traversal(root):
    # — BASE CASE —
    # Reached a None node (went past a leaf).
    # Nothing to visit here → return empty list.
    if root is None:
        return []

    # — LEFT → ROOT → RIGHT —
    # 1) FIRST recurse into left subtree (get all smaller values)
    # 2) THEN add current node's data (the 'root' in L-R-R)
    # 3) THEN recurse into right subtree (get all larger values)
    #
    # Result: left_values + [current] + right_values
    # In a BST: this is always in ASCENDING ORDER!
    return (
        in_order_traversal(root.left) + # all smaller values first
        [root.data] + # then this node
        in_order_traversal(root.right) # then all larger values
    )
```

```
# Time Complexity: O(n) - must visit every node exactly once
# Space Complexity:
#   Call stack: O(h) - h = height (O(log n) balanced, O(n) skewed)
#   Result list: O(n) - stores all n values
# KEY PROPERTY: On a BST → always gives SORTED output!
```

C++

```
void inOrder(Node* root, vector<int>& result) {
    if (root == nullptr) return;

    inOrder(root->left, result);    // left FIRST
    result.push_back(root->data);    // then root
    inOrder(root->right, result);   // then right
}
```

8.3 Full Recursion Trace

```
in_order_traversal(Node(1)):
├─ in_order(Node(3)) + [1] + in_order(Node(5))
│
├─ in_order(Node(3)):
│   ├─ in_order(Node(2)) + [3] + in_order(Node(4))
│   │
│   ├─ in_order(Node(2)):
│   │   ├─ in_order(None) → []
│   │   ├─ [2]
│   │   ├─ in_order(None) → []
│   │   └─ returns [] + [2] + [] = [2]
│   │
│   └─ in_order(Node(4)):
│       └─ returns [4]
│   └─ returns [2] + [3] + [4] = [2, 3, 4]
├─ [1]
├─ in_order(Node(5)):
│   ├─ in_order(None) → []
│   ├─ [5]
│   └─ in_order(Node(8)) → [8]
│       └─ returns [5, 8]
└─ Final: [2,3,4] + [1] + [5,8] = [2, 3, 4, 1, 5, 8]
```

NOTE: This tree is NOT a BST (values aren't in BST order),
so inorder doesn't give perfect sort here.

On a proper BST (12,15,18,20,30,40), inorder = [12,15,18,20,30,40] ✓

3 9. Post-Order Traversal

Left → Right → Root

9.1 The Pattern

Visit LEFT subtree first, then RIGHT subtree, then the ROOT last. Children are always processed before their parent. This is the 'clean up after children' traversal.

```

Tree:                Postorder Result:

    [1]
   /  \
  [3]  [5]
 /  \  \
[2] [4] [8]

Go to leaves first:
visit 2 (leftmost leaf)
visit 4 (right leaf of 3)
visit 3 (after both children done)
visit 8 (only child of 5)
visit 5 (after child done)
visit 1 (root — always LAST)

Postorder: 2 → 4 → 3 → 8 → 5 → 1
Notice: Root (1) is ALWAYS last in postorder!

```

9.2 Code — Line by Line

Python

```

def post_order_traversal(root):
    # — BASE CASE —
    # Reached a None node → nothing here → return empty list
    if root is None:
        return []

    # — LEFT → RIGHT → ROOT —
    # 1) FIRST recurse into left subtree completely
    # 2) THEN  recurse into right subtree completely
    # 3) LAST  add the current root's value
    #
    # Why is this useful?
    # Parent is processed AFTER all its descendants.
    # Perfect for: deletion (delete children before parent),
    # dependency resolution (finish deps before task),
    # postfix expression evaluation
    return (
        post_order_traversal(root.left) + # left subtree first
        post_order_traversal(root.right) + # right subtree second
        [root.data]                        # current node LAST
    )

```

```
# Time: O(n) — every node visited exactly once
# Space: O(h) call stack + O(n) result list
# KEY: Root is ALWAYS the last element in postorder output
```

C++

```
void postOrder(Node* root, vector<int>& result) {
    if (root == nullptr) return;

    postOrder(root->left, result);    // left first
    postOrder(root->right, result);   // right second
    result.push_back(root->data);     // root LAST
}
```

9.3 Full Recursion Trace

```
post_order_traversal(Node(1)):
├─ post_order(Node(3)) + post_order(Node(5)) + [1]
├─ post_order(Node(3)):
│   ├─ post_order(Node(2)) + post_order(Node(4)) + [3]
│   │   └─ post_order(Node(2)):
│   │       └─ post_order(None) → []
│   │       └─ post_order(None) → []
│   │       └─ returns [] + [] + [2] = [2]
│   │   └─ post_order(Node(4)):
│   │       └─ returns [4]
│   │   └─ returns [2] + [4] + [3] = [2, 4, 3]
├─ post_order(Node(5)):
│   └─ post_order(Node(8)) → [8]
│   └─ returns [] + [8] + [5] = [8, 5]
└─ Final: [2,4,3] + [8,5] + [1] = [2, 4, 3, 8, 5, 1] ✓

Root 1 appears LAST ✓
Both children of any node appear before it ✓
```

□ Postorder Use Cases:

Tree deletion: always delete children before parent (no orphaned nodes)

Postfix (RPN) expression evaluation: operands before operator

Folder size calculation: calculate subfolder size before parent folder

Dependency resolution: finish dependencies before the dependent task

Build systems: compile dependencies before the module that uses them

10. Complete Program

All operations together — Python & C++

10.1 Full Python Program

```
# =====
# COMPLETE BST IMPLEMENTATION
# Insertion | Search | Deletion | Traversals
# =====

class Node:
    def __init__(self, value):
        self.data = value
        self.left = None
        self.right = None

# — INSERTION —
# Time: Best O(1), Avg O(log n), Worst O(n) | Space: O(h)
def insertion_in_bst(root, value):
    if root is None:
        return Node(value)
    elif value == root.data:
        return root
    elif value < root.data:
        root.left = insertion_in_bst(root.left, value)
    else:
        root.right = insertion_in_bst(root.right, value)
    return root

# — SEARCH —
# Time: Best O(1), Avg O(log n), Worst O(n) | Space: O(h)
def search_in_bst(root, value):
    if root is None:
        print(f'{value} not found'); return False
    elif value == root.data:
        print(f'{value} found'); return True
    elif value < root.data:
        return search_in_bst(root.left, value)
    else:
        return search_in_bst(root.right, value)

# — INORDER SUCCESSOR HELPER —
# Time: O(h) | Space: O(1)
def get_in_order_successor(root):
    root = root.right
    while root is not None and root.left is not None:
        root = root.left
    return root

# — DELETION —
# Time: Best O(1), Avg O(log n), Worst O(n) | Space: O(h)
def delete_from_bst(root, value):
    if root is None:
        print(f'{value} not found'); return None
    elif value < root.data:
        root.left = delete_from_bst(root.left, value)
    elif value > root.data:
        root.right = delete_from_bst(root.right, value)
```

```

    else:
        if root.left is None: return root.right # Case 1/2
        if root.right is None: return root.left # Case 2
        succ = get_in_order_successor(root) # Case 3
        root.data = succ.data
        root.right = delete_from_bst(root.right, succ.data)
    return root

# — TRAVERSALS —————
# All: Time O(n), Space O(h) call stack + O(n) result
def pre_order_traversal(root):
    if root is None: return []
    return [root.data] + pre_order_traversal(root.left) +
pre_order_traversal(root.right)

def in_order_traversal(root):
    if root is None: return []
    return in_order_traversal(root.left) + [root.data] +
in_order_traversal(root.right)

def post_order_traversal(root):
    if root is None: return []
    return post_order_traversal(root.left) + post_order_traversal(root.right) +
[root.data]

# — BUILD BST —————
root = None
for val in [20, 15, 30, 12, 18, 40, 25, 50]:
    root = insertion_in_bst(root, val)

# — OUTPUTS —————
print('Inorder (sorted):', in_order_traversal(root))
# → [12, 15, 18, 20, 25, 30, 40, 50]
print('Preorder      :', pre_order_traversal(root))
# → [20, 15, 12, 18, 30, 25, 40, 50]
print('Postorder     :', post_order_traversal(root))
# → [12, 18, 15, 25, 50, 40, 30, 20]

search_in_bst(root, 18) # 18 found
search_in_bst(root, 100) # 100 not found

root = delete_from_bst(root, 12) # delete leaf
print('After deleting 12:', in_order_traversal(root))
# → [15, 18, 20, 25, 30, 40, 50]

root = delete_from_bst(root, 30) # delete 2-child node
print('After deleting 30:', in_order_traversal(root))
# → [15, 18, 20, 25, 40, 50]

```

10.2 Full C++ Program

```
#include <iostream>
#include <vector>
using namespace std;

struct Node {
    int data;
    Node* left;
    Node* right;
    Node(int v) : data(v), left(nullptr), right(nullptr) {}
};

Node* insert(Node* root, int val) {
    if (!root) return new Node(val);
    if (val < root->data) root->left = insert(root->left, val);
    else if (val > root->data) root->right = insert(root->right, val);
    return root;
}

bool search(Node* root, int val) {
    if (!root) return false;
    if (val == root->data) return true;
    return val < root->data ? search(root->left, val) : search(root->right,
val);
}

Node* getSuccessor(Node* root) {
    root = root->right;
    while (root && root->left) root = root->left;
    return root;
}

Node* deleteNode(Node* root, int val) {
    if (!root) return nullptr;
    if (val < root->data) { root->left = deleteNode(root->left, val); }
    else if (val > root->data) { root->right = deleteNode(root->right, val); }
    else {
        if (!root->left) { Node* t = root->right; delete root; return t; }
        if (!root->right) { Node* t = root->left; delete root; return t; }
        Node* succ = getSuccessor(root);
        root->data = succ->data;
        root->right = deleteNode(root->right, succ->data);
    }
    return root;
}

void preOrder(Node* root, vector<int>& v) {
    if (!root) return;
    v.push_back(root->data);
    preOrder(root->left, v);
```

```
        preOrder(root->right, v);
    }

    void inOrder(Node* root, vector<int>& v) {
        if (!root) return;
        inOrder(root->left, v);
        v.push_back(root->data);
        inOrder(root->right, v);
    }

    void postOrder(Node* root, vector<int>& v) {
        if (!root) return;
        postOrder(root->left, v);
        postOrder(root->right, v);
        v.push_back(root->data);
    }

    int main() {
        Node* root = nullptr;
        for (int v : {20, 15, 30, 12, 18, 40, 25, 50})
            root = insert(root, v);

        vector<int> res;
        inOrder(root, res);
        // prints: 12 15 18 20 25 30 40 50
        return 0;
    }
```



11. Master Complexity Reference

Every operation, every case

11.1 BST Operations

Operation	Best Case	Average Case	Worst Case	Space	When worst occurs
Insert	$O(1)$	$O(\log n)$	$O(n)$	$O(h)$	Sorted input → skewed tree
Search	$O(1)$	$O(\log n)$	$O(n)$	$O(h)$	Element at deepest leaf or absent
Delete (leaf)	$O(\log n)$	$O(\log n)$	$O(n)$	$O(h)$	Leaf at deepest level of skewed tree
Delete (2 child)	$O(\log n)$	$O(\log n)$	$O(n)$	$O(h)$	Skewed tree + node near bottom
Find Min	$O(1)$	$O(\log n)$	$O(n)$	$O(1)$	Leftmost node in skewed-right tree
Find Max	$O(1)$	$O(\log n)$	$O(n)$	$O(1)$	Rightmost node in skewed-left tree
Get Successor	$O(1)$	$O(\log n)$	$O(n)$	$O(1)$	Successor at deepest level

11.2 Traversal Complexities

Traversal	Pattern	Time	Stack Space (balanced)	Stack Space (skewed)	Result Space	Key Use
Preorder	Root→L→R	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	Copy tree, serialize
Inorder	L→Root→R	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	Sorted output from BST
Postorder	L→R→Root	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	Delete tree, postfix eval

11.3 Key Rules to Memorize

- Inorder of a BST = sorted array. This is THE most important BST property.
- Preorder first element = root. Postorder last element = root. Inorder = root in middle.
- To reconstruct a tree: you need preorder + inorder OR postorder + inorder (not just one).
- All traversals are $O(n)$ time — you must visit every node exactly once.
- Traversal space = $O(h)$: balanced $O(\log n)$, skewed $O(n)$ (recursion stack depth).
- BST worst case is ALWAYS when input is sorted → creates a right-skewed tree.
- Use AVL or Red-Black trees in production to prevent worst case $O(n)$.

□ The 3 Interview Questions That Are The Same:

Q1: Print BST elements in increasing order → INORDER traversal

Q2: Check if a tree is a valid BST → use inorder, check if result is sorted

Q3: Inorder traversal of BST → Left → Root → Right

All 3 answers = the same code. Know this and you get all 3 for free!

Master the recursion. Master the tree. □

Every BST operation is just recursion + the BST rule applied at each node.

Understand the recursion trace, and no BST interview question can surprise you.